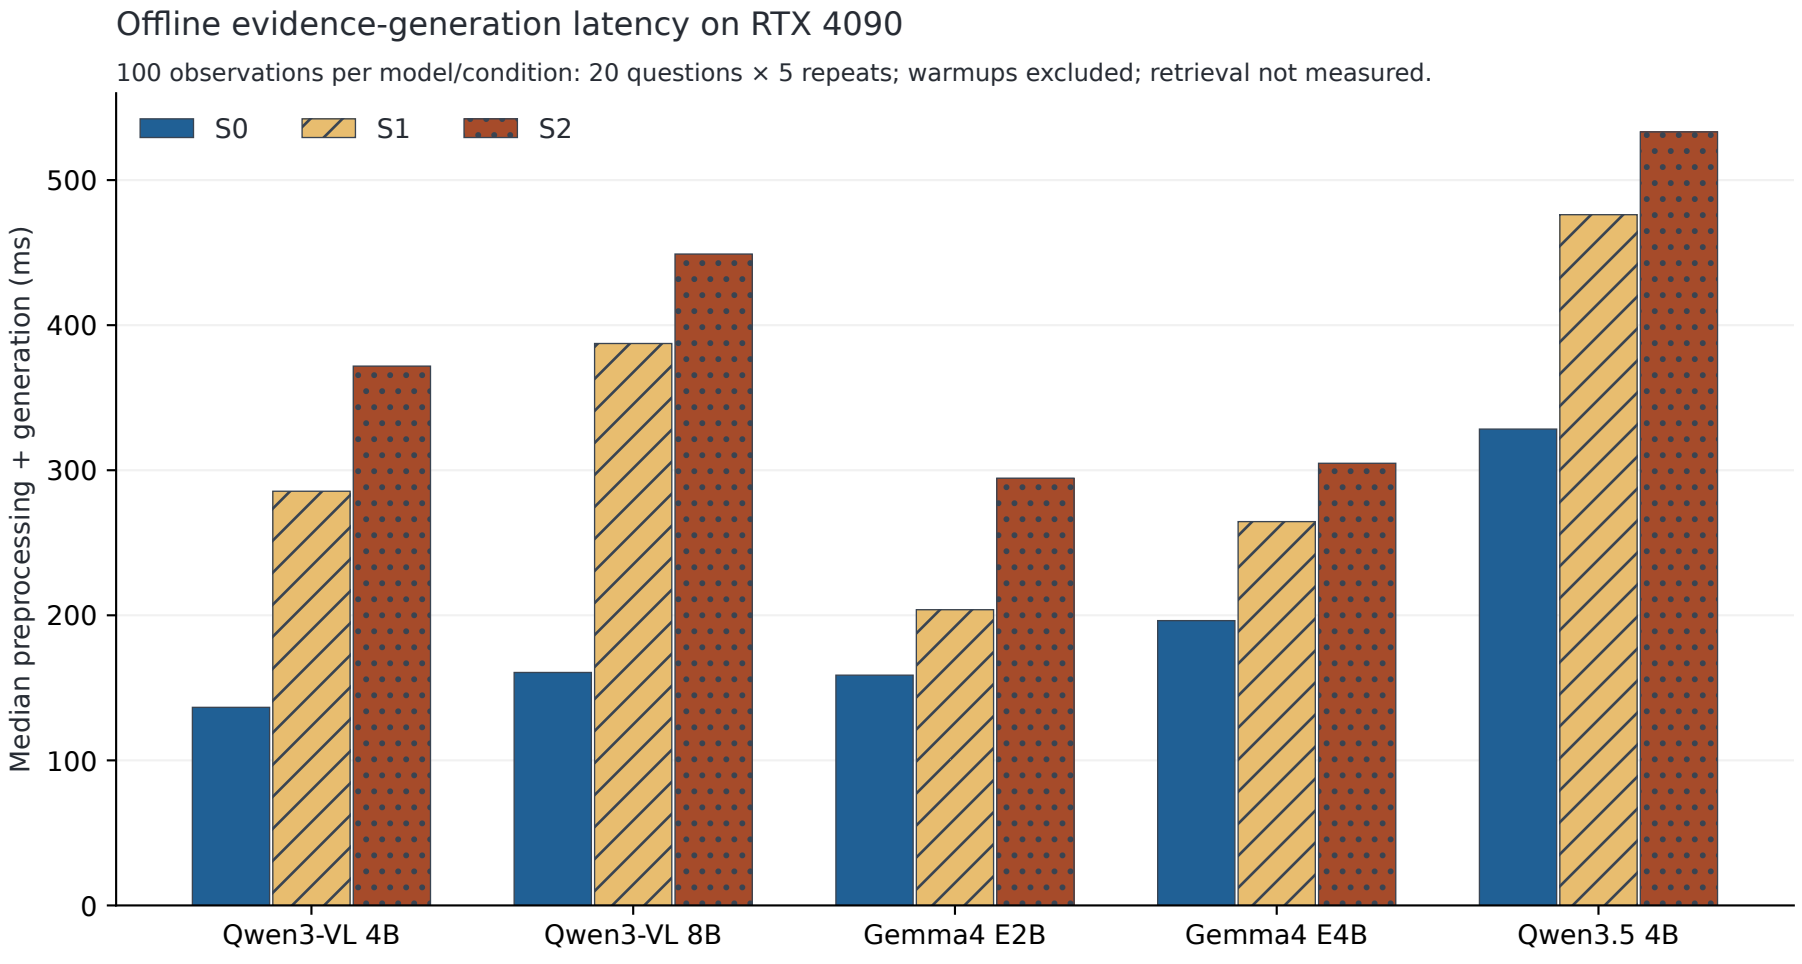
Offline evidence-generation latency on RTX 4090
100 observations per model/condition: 20 questions × 5 repeats; warmups excluded; retrieval not measured.
S0
S1
S2
Median preprocessing + generation (ms)
0
100
200
300
400
500
Qwen3-VL 4B
Qwen3-VL 8B
Gemma4 E2B
Gemma4 E4B
Qwen3.5 4B